

# pm2-2-quality-prediction-in-a-mining-process-filled

February 15, 2022

## 1 Quality Prediction in a Mining Process by using RNN

In this notebook, we are going to predict how much impurity is in the ore concentrate. As this impurity is measured every hour, if we can predict how much silica (impurity) is in the ore concentrate, we can help the engineers, giving them early information to take actions. Hence, they will be able to take corrective actions in advance (reduce impurity, if it is the case) and also help the environment (reducing the amount of ore that goes to tailings as you reduce silica in the ore concentrate). To this end, we are going to use the dataset **Quality Prediction in a Mining Process Data** from [Kaggle](#).

In order to have a clean notebook, some functions are implemented in the file *utils.py* (e.g., `plot_loss_and_accuracy`).

Summary: - Data Pre-processing - Data Visualisation - Data Normalisation - Building the Models - Splitting the Data into Train and Test Sets - Gated Recurrent Unit (GRU) - Long-short Term Memory (LSTM)

**All the libraries used in this notebook are Open Source.**

```
[1]: # Standard libraries - no deep learning yet
import numpy as np # written in C, is faster and robust library for numerical,
    ↪ and matrix operations
import pandas as pd # data manipulation library, it is widely used for data,
    ↪ analysis and relies on numpy library.
import matplotlib.pyplot as plt # for plotting
from datetime import datetime # supplies classes for manipulating dates and,
    ↪ times in both simple and complex ways

from utils import *

# the following to lines will tell to the python kernel to always update the,
    ↪ kernel for every utils.py
# modification, without the need of restarting the kernel.
# Of course, for every modification in util.py, we need to reload this cell
%load_ext autoreload
%autoreload 2
%matplotlib inline
```

Using TensorFlow backend.

## 1.1 Data Pre-processing

First download the dataset (click [here](#)) unzip `quality-prediction-in-a-mining-process.zip`

The **Quality Prediction in a Mining Process Data** includes ([Kaggle](#)):

- The first column shows time and date range (from march of 2017 until september of 2017). Some columns were sampled every 20 second. Others were sampled on a hourly base. *This make the data processing harder, however, for this tutorial we will not re-sample the data.*
- The second and third columns are quality measures of the iron ore pulp right before it is fed into the flotation plant.
- Column 4 until column 8 are the most important variables that impact in the ore quality in the end of the process.
- Column 9 until column 22, we can see process data level and air flow inside the flotation columns, which also impact in ore quality.
- The last two columns are the final iron ore pulp quality measurement from the lab. Target is to predict the last column, which is the % of silica in the iron ore concentrate.

We are going to use [Pandas](#) for the data processing. The function `read_csv` is going to be used to read the csv file.

```
[2]: dataset = pd.read_csv('../data/MiningProcess_Flotation_Plant_Database.
    ↪ csv', index_col=0, decimal=",")

# Set the index name to 'date'
dataset.index.name = 'date'

dataset.head()
```

```
[2]:
```

	% Iron Feed	% Silica Feed	Starch Flow	Amina Flow \
date				
2017-03-10 01:00:00	55.2	16.98	3019.53	557.434
2017-03-10 01:00:00	55.2	16.98	3024.41	563.965
2017-03-10 01:00:00	55.2	16.98	3043.46	568.054
2017-03-10 01:00:00	55.2	16.98	3047.36	568.665
2017-03-10 01:00:00	55.2	16.98	3033.69	558.167

	Ore Pulp Flow	Ore Pulp pH	Ore Pulp Density \
date			
2017-03-10 01:00:00	395.713	10.0664	1.74
2017-03-10 01:00:00	397.383	10.0672	1.74
2017-03-10 01:00:00	399.668	10.0680	1.74
2017-03-10 01:00:00	397.939	10.0689	1.74
2017-03-10 01:00:00	400.254	10.0697	1.74

	Flotation Column 01 Air Flow \
date	
2017-03-10 01:00:00	249.214

2017-03-10 01:00:00	249.719
2017-03-10 01:00:00	249.741
2017-03-10 01:00:00	249.917
2017-03-10 01:00:00	250.203

#### Flotation Column 02 Air Flow \

date	
2017-03-10 01:00:00	253.235
2017-03-10 01:00:00	250.532
2017-03-10 01:00:00	247.874
2017-03-10 01:00:00	254.487
2017-03-10 01:00:00	252.136

#### Flotation Column 03 Air Flow ... \

date		
2017-03-10 01:00:00	250.576	...
2017-03-10 01:00:00	250.862	...
2017-03-10 01:00:00	250.313	...
2017-03-10 01:00:00	250.049	...
2017-03-10 01:00:00	249.895	...

#### Flotation Column 07 Air Flow Flotation Column 01 Level \

date		
2017-03-10 01:00:00	250.884	457.396
2017-03-10 01:00:00	248.994	451.891
2017-03-10 01:00:00	248.071	451.240
2017-03-10 01:00:00	251.147	452.441
2017-03-10 01:00:00	248.928	452.441

#### Flotation Column 02 Level Flotation Column 03 Level \

date		
2017-03-10 01:00:00	432.962	424.954
2017-03-10 01:00:00	429.560	432.939
2017-03-10 01:00:00	468.927	434.610
2017-03-10 01:00:00	458.165	442.865
2017-03-10 01:00:00	452.900	450.523

#### Flotation Column 04 Level Flotation Column 05 Level \

date		
2017-03-10 01:00:00	443.558	502.255
2017-03-10 01:00:00	448.086	496.363
2017-03-10 01:00:00	449.688	484.411
2017-03-10 01:00:00	446.210	471.411
2017-03-10 01:00:00	453.670	462.598

#### Flotation Column 06 Level Flotation Column 07 Level \

date

2017-03-10 01:00:00	446.370	523.344
2017-03-10 01:00:00	445.922	498.075
2017-03-10 01:00:00	447.826	458.567
2017-03-10 01:00:00	437.690	427.669
2017-03-10 01:00:00	443.682	425.679

	% Iron Concentrate	% Silica Concentrate
date		
2017-03-10 01:00:00	66.91	1.31
2017-03-10 01:00:00	66.91	1.31
2017-03-10 01:00:00	66.91	1.31
2017-03-10 01:00:00	66.91	1.31
2017-03-10 01:00:00	66.91	1.31

[5 rows x 23 columns]

Given our time and computational resources restrictions, we are going to select the first 100,000 observations for this tutorial.


```
[3]: dataset = dataset.iloc[:100000,:]
```

### 1.1.1 Data Visualisation

```
[4]: # Plotting the Silica Concentrate
plt.figure(figsize = (15, 5))
plt.xlabel("x")
plt.ylabel("Silica Concentrate")

dataset['% Silica Concentrate'].plot()
```

```
[4]: <AxesSubplot:xlabel='date', ylabel='Silica Concentrate'>
```



## 1.2 Data Normalisation

Here are going to normalise all the features and transform the data into a supervised learning problem. The features to be predicted are removed, as we would like to predict just the *Silica Concentrate* (last element in every feature array).

### 1.2.1 Transforming the data into a supervised learning problem

This step will involve framing the dataset as a **supervised learning problem**. As we would like to predict the “silica concentrate”, we will set the corresponding column to be the output (label  $y$ ).

We would like to predict the silica concentrate ( $y_t$ ) at the current time ( $t$ ) given the measurements at the prior time steps (lets say  $t-1, t-2, \dots, t-n$ , in which  $n$  is the number of past observations to be used to forecast  $y_t$ ).

The function `create_window` (see `utils.py`) converts the time-series to a supervised learning problem. The new dataset is constructed as a **DataFrame**, with each column suitably named both by variable number and time step, for example, `var1(t-1)` for **%Iron Feed** at the previous observation ( $t-1$ ). This allows you to design a variety of different time step sequence type forecasting problems from a given univariate or multivariate time series.

```
[4]: # Scikit learn libraries
from sklearn.preprocessing import MinMaxScaler #Allows normalisation

# Convert the data to float
values = dataset.values.astype('float32')

# Normalise features
scaler = MinMaxScaler(feature_range=(0, 1))
scaled = scaler.fit_transform(values)

# Specify the number of lag
n_in = 5
n_features = 23

# Transform the time-series to a supervised learning problem representation
reframed = create_window(scaled, n_in = n_in, n_out = 1, drop_nan = True)

# Summarise the new frames (reframes)
print(reframed.head(1))
```

```
   var1(t-5)  var2(t-5)  var3(t-5)  var4(t-5)  var5(t-5)  var6(t-5)  \
5    0.476715    0.502299    0.483124    0.665398    0.459145    0.641907

   var7(t-5)  var8(t-5)  var9(t-5)  var10(t-5)  ...  var14(t)  var15(t)  \
5    0.660635    0.377486    0.432691      0.4025  ...    0.382131    0.41283

   var16(t)  var17(t)  var18(t)  var19(t)  var20(t)  var21(t)  var22(t)  \
5    0.375913    0.4394    0.535916    0.559504    0.511396    0.516056    0.875677
```

```

    var23(t)
5  0.114165

[1 rows x 138 columns]

```

### 1.3 Building the Models

So far, we just preprocessed the dataset. Now, we are going to build the following sequential models:

- Gated Recurrent Unit (GRU)
- Long Short-Term Memory (LSTM)

The models consists in a **many\_to\_one** architecture, in which the input is a **sequence** of the past observations and the output is the predicted value (in this case with dimension equal 1).

```

[5]: from keras.models import Sequential
      from keras.layers import Dense, Dropout
      from keras.layers import LSTM, GRU

      from sklearn.metrics import mean_squared_error # allows compute the mean square
      →error to performance analysis

```

#### 1.3.1 Splitting the Data into Train and Test Sets

```

[6]: # split into train and test sets
      values = reframed.values

      # We will use 80% of the data for training and 20% for testing
      n_train = round(0.8 * dataset.shape[0])
      train = values[:n_train, :]
      test = values[n_train:, :]

      # Split into input and outputs
      n_obs = n_in * n_features # the number of total features is given by the number
      →of past
                                   # observations * number of features. In this case we
      →have
                                   # 5 past observations and 23 features, so the number
      →of total
                                   # features is 115.
      x_train, y_train = train[:, :n_obs], train[:, n_features-1] # note that fore
      →y_train, we are removing
                                   # just the last
      →observation of the
                                   # silica
      →concentrate

      x_test, y_test = test[:, :n_obs], test[:, n_features-1]
      print('Number of total features (n_obs): ', x_train.shape[1])

```

```

print('Number of samples in training set: ', x_train.shape[0])
print('Number of samples in testing set: ', x_test.shape[0])

# Reshape input to be 3D [samples, timesteps, features]
x_train = x_train.reshape((x_train.shape[0], n_in, n_features))
x_test = x_test.reshape((x_test.shape[0], n_in, n_features))

```

```

Number of total features (n_obs): 115
Number of samples in training set: 80000
Number of samples in testing set: 19995

```

### 1.3.2 Gated Recurrent Unit (GRU)

To build the model, we are going to use the following components from Keras:

- [Sequential](#): allows us to create models layer-by-layer.
- [GRU](#): provides a GRU architecture
- [Dense](#): provides a regular fully-connected layer
- [Activation](#): defines the activation function to be used

Basically, we can define the sequence of the model by using *Sequential()*:

```

model = Sequential()
model.add(GRU(...))
...

```

where the function *add(...)* stacks the layers. Once created the model, we can configure the training by using the function [compile](#). Here we need to define the [loss](#) function (mean squared error, mean absolute error, cosine proximity, among others.) and the [optimizer](#) (Stochastic gradient descent, RMSprop, adam, among others), as follows:

```

model.compile(loss = "...",
              optimizer = "...")

```

Also, we have the option to see a summary representation of the model by using the function [summary](#). This function summarises the model and tells us the number of parameters that we need to tune.

LLPS-2022-02-03:

- the `model_gru.add(GRU...` below raises `*"NotImplementedError: Cannot convert a symbolic Tensor (gru_2/strided_slice:0) to a numpy array. This error may indicate that you're trying to pass a Tensor to a NumPy call, which is not supported"`
- problem seems to be related to numpy version in this Docker image `rio_cs_p1 hari-doop/cs_p1:latest`, which is '1.21.2'. LLPS `np.__version__` is '1.20.3' and works!
- People report that TF was built with 'numpy ~ 1.19.2' so, that's the reason! REF: <https://github.com/tensorflow/tensorflow/issues/47242>

```
[7]: np.__version__
```

```
[7]: '1.20.3'
```

```
[8]: # Define the model.
model_gru = Sequential()

# the input_shape is the number of past observations (n_in) and the number of
# features
# per past observations (23)
model_gru.add(GRU(input_shape=(x_train.shape[1], x_train.shape[2]),
                    units = 128,
                    return_sequences = False))

model_gru.add(Dense(units=1))

# We compile the model by defining the mean absolute error (denoted by mae) as
# loss function and
# adam as optimizer
model_gru.compile(loss = "mae",
                  optimizer = "adam")

# just print the model
model_gru.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
gru_1 (GRU)	(None, 128)	58368
dense_1 (Dense)	(None, 1)	129

Total params: 58,497  
 Trainable params: 58,497  
 Non-trainable params: 0

### 1.3.3 Training the Model

Once defined the model, we need to train it by using the function `fit`. This function performs the optimisation step. Hence, we can define the following parameters such as:

- batch size: defines the number of samples that will be propagated through the network
- epochs: defines the number of times in which all the training set (`x_train_scaled`) are used once to update the weights
- validation split: defines the percentage of training data to be used for validation
- among others (click [here](#) for more information)

This function return the *history* of the training, that can be used for further performance analysis.



```
[9]: # Training
hist_gru = model_gru.fit(x_train, y_train,
                        epochs=50,
                        batch_size=256,
                        validation_split = 0.1,
                        verbose=1, # To not print the output, set verbose=0
                        shuffle=False)
```

WARNING:tensorflow:From /home/leo/anaconda3/envs/day09/lib/python3.7/site-packages/keras/backend/tensorflow\_backend.py:422: The name tf.global\_variables is deprecated. Please use tf.compat.v1.global\_variables instead.

2022-02-15 19:09:36.141701: I tensorflow/core/platform/cpu\_feature\_guard.cc:142] Your CPU supports instructions that this TensorFlow binary was not compiled to use: SSE4.1 SSE4.2 AVX AVX2 AVX512F FMA

2022-02-15 19:09:36.179091: I tensorflow/core/platform/profile\_utils/cpu\_utils.cc:94] CPU Frequency: 2599990000 Hz

2022-02-15 19:09:36.180794: I tensorflow/compiler/xla/service/service.cc:168]

XLA service 0x55d56f6c9980 executing computations on platform Host. Devices:

2022-02-15 19:09:36.180829: I tensorflow/compiler/xla/service/service.cc:175] StreamExecutor device (0): <undefined>, <undefined>

2022-02-15 19:09:36.240908: W

tensorflow/compiler/jit/mark\_for\_compilation\_pass.cc:1412] (One-time warning): Not using XLA:CPU for cluster because envvar

TF\_XLA\_FLAGS=--tf\_xla\_cpu\_global\_jit was not set. If you want XLA:CPU, either set that envvar, or use experimental\_jit\_scope to enable XLA:CPU. To confirm that XLA is active, pass --vmodule=xla\_compilation\_cache=1 (as a proper command-line flag, not via TF\_XLA\_FLAGS) or set the envvar XLA\_FLAGS=--xla\_hlo\_profile.

Train on 72000 samples, validate on 8000 samples

Epoch 1/50

72000/72000 [=====] - 6s 89us/step - loss: 0.1273 - val\_loss: 0.0462

Epoch 2/50

72000/72000 [=====] - 6s 86us/step - loss: 0.0531 - val\_loss: 0.0142

Epoch 3/50

72000/72000 [=====] - 6s 85us/step - loss: 0.0584 - val\_loss: 0.0571

Epoch 4/50

72000/72000 [=====] - 6s 83us/step - loss: 0.0360 - val\_loss: 0.0213

Epoch 5/50

72000/72000 [=====] - 6s 78us/step - loss: 0.0320 - val\_loss: 0.0187

Epoch 6/50

```

72000/72000 [=====] - 6s 84us/step - loss: 0.0259 -
val_loss: 0.0139
Epoch 7/50
72000/72000 [=====] - 6s 81us/step - loss: 0.0254 -
val_loss: 0.0358
Epoch 8/50
72000/72000 [=====] - 6s 81us/step - loss: 0.0236 -
val_loss: 0.0098
Epoch 9/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0151 -
val_loss: 0.0097
Epoch 10/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0186 -
val_loss: 0.0148
Epoch 11/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0227 -
val_loss: 0.0250
Epoch 12/50
72000/72000 [=====] - 5s 75us/step - loss: 0.0248 -
val_loss: 0.0177
Epoch 13/50
72000/72000 [=====] - 6s 82us/step - loss: 0.0146 -
val_loss: 0.0128
Epoch 14/50
72000/72000 [=====] - 5s 72us/step - loss: 0.0231 -
val_loss: 0.0209
Epoch 15/50
72000/72000 [=====] - 6s 78us/step - loss: 0.0205 -
val_loss: 0.0142
Epoch 16/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0198 -
val_loss: 0.0126
Epoch 17/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0150 -
val_loss: 0.0128
Epoch 18/50
72000/72000 [=====] - 6s 82us/step - loss: 0.0237 -
val_loss: 0.0143
Epoch 19/50
72000/72000 [=====] - 6s 79us/step - loss: 0.0205 -
val_loss: 0.0170
Epoch 20/50
72000/72000 [=====] - 6s 78us/step - loss: 0.0211 -
val_loss: 0.0058
Epoch 21/50
72000/72000 [=====] - 6s 77us/step - loss: 0.0146 -
val_loss: 0.0312
Epoch 22/50

```

```

72000/72000 [=====] - 6s 83us/step - loss: 0.0186 -
val_loss: 0.0133
Epoch 23/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0164 -
val_loss: 0.0108
Epoch 24/50
72000/72000 [=====] - 6s 84us/step - loss: 0.0107 -
val_loss: 0.0105
Epoch 25/50
72000/72000 [=====] - 6s 82us/step - loss: 0.0172 -
val_loss: 0.0103
Epoch 26/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0158 -
val_loss: 0.0075
Epoch 27/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0115 -
val_loss: 0.0068
Epoch 28/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0146 -
val_loss: 0.0098
Epoch 29/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0171 -
val_loss: 0.0212
Epoch 30/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0129 -
val_loss: 0.0087
Epoch 31/50
72000/72000 [=====] - 6s 84us/step - loss: 0.0146 -
val_loss: 0.0136
Epoch 32/50
72000/72000 [=====] - 6s 77us/step - loss: 0.0188 -
val_loss: 0.0317
Epoch 33/50
72000/72000 [=====] - 5s 75us/step - loss: 0.0157 -
val_loss: 0.0080
Epoch 34/50
72000/72000 [=====] - 6s 77us/step - loss: 0.0174 -
val_loss: 0.0114
Epoch 35/50
72000/72000 [=====] - 6s 84us/step - loss: 0.0153 -
val_loss: 0.0077
Epoch 36/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0095 -
val_loss: 0.0159
Epoch 37/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0132 -
val_loss: 0.0158
Epoch 38/50

```

```


72000/72000 [=====] - 6s 86us/step - loss: 0.0201 -
val_loss: 0.0440
Epoch 39/50
72000/72000 [=====] - 6s 83us/step - loss: 0.0138 -
val_loss: 0.0069
Epoch 40/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0132 -
val_loss: 0.0273
Epoch 41/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0201 -
val_loss: 0.0163
Epoch 42/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0119 -
val_loss: 0.0062
Epoch 43/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0147 -
val_loss: 0.0054
Epoch 44/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0121 -
val_loss: 0.0051
Epoch 45/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0102 -
val_loss: 0.0141
Epoch 46/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0092 -
val_loss: 0.0045
Epoch 47/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0155 -
val_loss: 0.0102
Epoch 48/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0147 -
val_loss: 0.0093
Epoch 49/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0094 -
val_loss: 0.0115
Epoch 50/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0134 -
val_loss: 0.0091

```

### 1.3.4 Prediction and Performance Analysis

Here we can see if the model overfits or underfits. First, we are going to plot the 'loss' and the 'Accuracy' in from the training step.

```
[10]: plot_loss(hist_gru)
```



Once the model was trained, we can use the function `predict` for prediction tasks. We are going to use the function `inverse_transform` (see `utils.py`) to invert the scaling (transform the values to the original ones).

Given the predictions and expected values in their original scale, we can then compute the error score for the model.

```
[11]: yhat_gru = model_gru.predict(x_test)

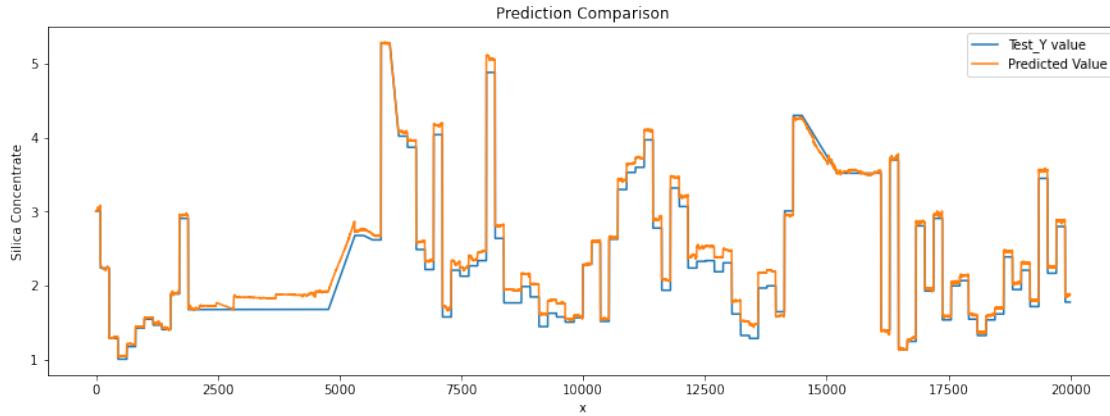
# performing the inverse transform on test_X and yhat_rnn
inv_y_gru, inv_yhat_gru = inverse_transform_multiple(x_test, y_test, yhat_gru,
↪ scaler, n_in, n_features)

# calculate RMSE
mse_gru = mean_squared_error(inv_y_gru, inv_yhat_gru)
print('Test MSE: %.3f' % mse_gru)
```


Test MSE: 0.023

### 1.3.5 Visualising the predicted Data


```
[12]: plot_comparison([inv_y_gru, inv_yhat_gru],
                      ['Test_Y value', 'Predicted Value'],
                      title='Prediction Comparison')
```



```
[13]: plot_comparison([inv_y_gru[0:300], inv_yhat_gru[0:300]],
                      ['Test_Y value', 'Predicted Value'],
                      title='Prediction Comparison of the first 300 Observations')
```



```
[14]: plot_comparison([inv_y_gru[5500:5800], inv_yhat_gru[5500:5800]],
                      ['Test_Y value', 'Predicted Value'],
                      title='Prediction Comparison of a Farthest 300 Observations')
```



### 1.3.6 Long-Short Term Memory (LSTM)

Now **you** are going to build the model based on LSTM. Like GRU, we are going to use the following components from Keras:

- **Sequential**: allows us to create models layer-by-layer.
- **LSTM**: provides a LSTM architecture
- **Dense**: provides a regular fully-connected layer
- **Activation**: defines the activation function to be used

Basically, you are going to define the sequence of the model by using *Sequential()*:

```
model = Sequential()
model.add(LSTM(...))
...
```

and configure the training by using the function *compile*:

```
model.compile(loss = "...",
              optimizer = "...")
```

Follow the below steps for this task.

**Step 1:** Create the model: 1) Define the number of layers (we suggest at this stage to use just one, but it is up to you) 2) Create the fully connected layer

For example:

```
# Define the model.
model_lstm = Sequential()

# Stacking just one LSTM
model_lstm.add(LSTM(input_shape=(train_X.shape[1], train_X.shape[2]),
                    units = 128,
                    return_sequences = False))
```

```
# Fully connected layer
model_lstm.add(Dense(units=1))
```

```
[15]: # Define the model.
model_lstm = Sequential()

# Stacking just one LSTM
model_lstm.add(LSTM(input_shape=(x_train.shape[1], x_train.shape[2]),
                      units = 128,
                      return_sequences = False))

# Fully connected layer
model_lstm.add(Dense(units=1))
```

**Step 2:** Configure the training: 1) Define the loss function (e.g., ‘mae’ for mean average error or ‘mse’ for mean squared error) 2) Define the optimiser (e.g., ‘adam’, ‘rmsprop’, ‘sgd’, ‘adagrad, etc)

For example:

```
model_lstm.compile(loss = "mae",
                  optimizer = "adam")
```

```
[16]: model_lstm.compile(loss = "mae",
                        optimizer = "adam")
```

**Step 3:** Call the function

```
model_lstm.summary()
```

to summarise the model.

```
[17]: model_lstm.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 128)	77824
dense_2 (Dense)	(None, 1)	129

Total params: 77,953  
 Trainable params: 77,953  
 Non-trainable params: 0

**Step 4:** Defined the number of epochs, validation\_split and batch\_size that best fit for you model and call the function fit to train the model.

For example:



```
hist_lstm = model_lstm.fit(train_X, train_y,
                           epochs=50,
                           batch_size=256,
                           validation_split = 0.1,
                           verbose=1, shuffle=False)
```

```
[18]: hist_lstm = model_lstm.fit(x_train, y_train,
                                epochs=50,
                                batch_size=256,
                                validation_split = 0.1,
                                verbose=1, shuffle=False)
```

Train on 72000 samples, validate on 8000 samples

```
Epoch 1/50
72000/72000 [=====] - 6s 88us/step - loss: 0.1084 -
val_loss: 0.0537
Epoch 2/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0534 -
val_loss: 0.0266
Epoch 3/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0438 -
val_loss: 0.0317
Epoch 4/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0381 -
val_loss: 0.0272
Epoch 5/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0455 -
val_loss: 0.0314
Epoch 6/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0312 -
val_loss: 0.0148
Epoch 7/50
72000/72000 [=====] - 6s 88us/step - loss: 0.0211 -
val_loss: 0.0156
Epoch 8/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0280 -
val_loss: 0.0487
Epoch 9/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0281 -
val_loss: 0.0474
Epoch 10/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0340 -
val_loss: 0.0228
Epoch 11/50
72000/72000 [=====] - 6s 84us/step - loss: 0.0295 -
val_loss: 0.0262
Epoch 12/50
72000/72000 [=====] - 6s 78us/step - loss: 0.0263 -
```

```

val_loss: 0.0246
Epoch 13/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0292 -
val_loss: 0.0146
Epoch 14/50
72000/72000 [=====] - 6s 89us/step - loss: 0.0175 -
val_loss: 0.0229
Epoch 15/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0340 -
val_loss: 0.0109
Epoch 16/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0162 -
val_loss: 0.0080
Epoch 17/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0183 -
val_loss: 0.0315
Epoch 18/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0218 -
val_loss: 0.0335
Epoch 19/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0240 -
val_loss: 0.0131
Epoch 20/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0208 -
val_loss: 0.0124
Epoch 21/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0152 -
val_loss: 0.0178
Epoch 22/50
72000/72000 [=====] - 6s 84us/step - loss: 0.0247 -
val_loss: 0.0122
Epoch 23/50
72000/72000 [=====] - 6s 76us/step - loss: 0.0237 -
val_loss: 0.0106
Epoch 24/50
72000/72000 [=====] - 6s 81us/step - loss: 0.0229 -
val_loss: 0.0098
Epoch 25/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0146 -
val_loss: 0.0094
Epoch 26/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0114 -
val_loss: 0.0227
Epoch 27/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0187 -
val_loss: 0.0129
Epoch 28/50
72000/72000 [=====] - 6s 88us/step - loss: 0.0131 -

```

```

val_loss: 0.0104
Epoch 29/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0160 -
val_loss: 0.0175
Epoch 30/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0137 -
val_loss: 0.0083
Epoch 31/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0168 -
val_loss: 0.0063
Epoch 32/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0170 -
val_loss: 0.0195
Epoch 33/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0154 -
val_loss: 0.0115
Epoch 34/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0143 -
val_loss: 0.0078
Epoch 35/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0159 -
val_loss: 0.0169
Epoch 36/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0157 -
val_loss: 0.0139
Epoch 37/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0230 -
val_loss: 0.0299
Epoch 38/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0141 -
val_loss: 0.0111
Epoch 39/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0154 -
val_loss: 0.0081
Epoch 40/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0118 -
val_loss: 0.0187
Epoch 41/50
72000/72000 [=====] - 6s 84us/step - loss: 0.0139 -
val_loss: 0.0129
Epoch 42/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0137 -
val_loss: 0.0073
Epoch 43/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0184 -
val_loss: 0.0080
Epoch 44/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0147 -

```


```

val_loss: 0.0094
Epoch 45/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0139 -
val_loss: 0.0181
Epoch 46/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0182 -
val_loss: 0.0160
Epoch 47/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0140 -
val_loss: 0.0073
Epoch 48/50
72000/72000 [=====] - 6s 86us/step - loss: 0.0113 -
val_loss: 0.0085
Epoch 49/50
72000/72000 [=====] - 6s 87us/step - loss: 0.0138 -
val_loss: 0.0128
Epoch 50/50
72000/72000 [=====] - 6s 85us/step - loss: 0.0091 -
val_loss: 0.0062

```

**Using the history** Here we can see if the model overfits or underfits

```
[19]: plot_loss(hist_lstm)
```



```
[24]: yhat_lstm = model_lstm.predict(x_test)


# performing the inverse transform on test_X and yhat_rnn
inv_y_lstm, inv_yhat_lstm = inverse_transform_multiple(x_test, y_test,
→yhat_lstm, scaler, n_in, n_features)

# calculate RMSE
mse_lstm = mean_squared_error(inv_y_lstm, inv_yhat_lstm)
print('Test MSE: %.3f' % mse_lstm)
```

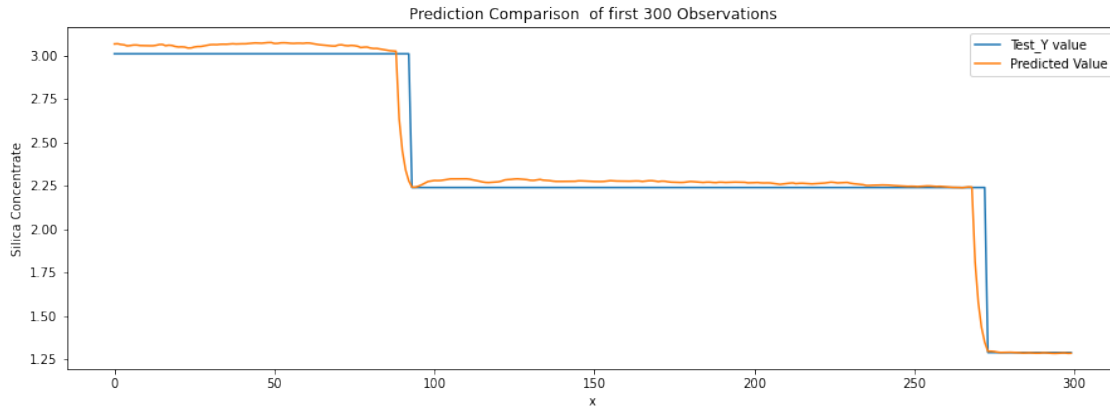
Test MSE: 0.012

### 1.3.7 Visualising the predicted Data


```
[25]: plot_comparison([inv_y_lstm, inv_yhat_lstm],
    ['Test_Y value', 'Predicted Value'],
    title='Prediction Comparison')
```



```
[26]: plot_comparison([inv_y_lstm[0:300], inv_yhat_lstm[0:300]],
    ['Test_Y value', 'Predicted Value'],
    title='Prediction Comparison of first 300 Observations')
```




```
[27]: plot_comparison([inv_y_lstm[5500:5800], inv_yhat_lstm[5500:5800]],
                      ['Test_Y value', 'Predicted Value'],
                      title='Prediction Comparison of a Farthest 300 Observations')
```



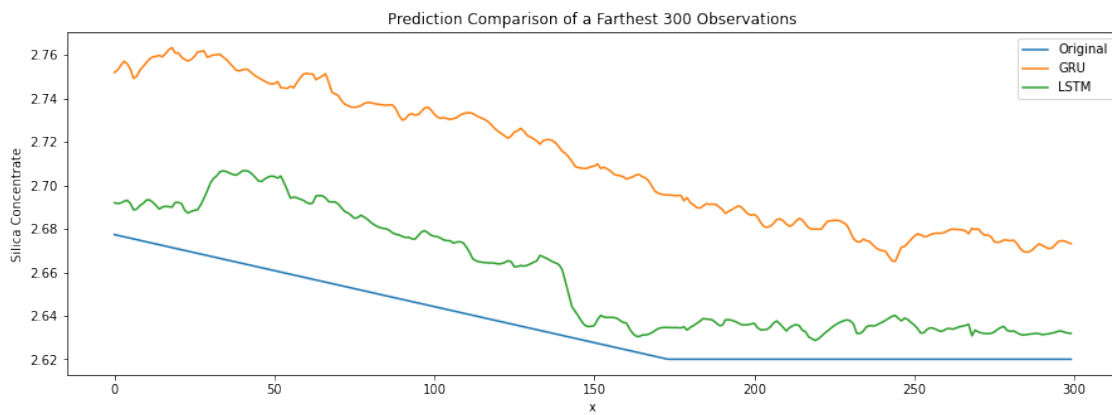
## 1.4 Models comparison

**Exercise:** run the code below and discuss the results.

```
[28]: plot_comparison([inv_y_lstm[0:300],
                      inv_yhat_gru[0:300], inv_yhat_lstm[0:300]],
                      ['Original', 'GRU', 'LSTM'],
                      title='Prediction Comparison of the First 300 Observations')
```



```
[29]: plot_comparison([inv_y_lstm[5500:5800],
                        inv_yhat_gru[5500:5800], inv_yhat_lstm[5500:5800]],
                        ['Original', 'GRU', 'LSTM'],
                        title='Prediction Comparison of a Farthest 300 Observations')
```



```
[30]: print('Comparing the MSE of the three models:')
print('  GRU: ', mse_gru)
print('  LSTM: ', mse_lstm)
```

Comparing the MSE of the three models:

GRU: 0.022612361

LSTM: 0.012126184

```
[ ]:
```